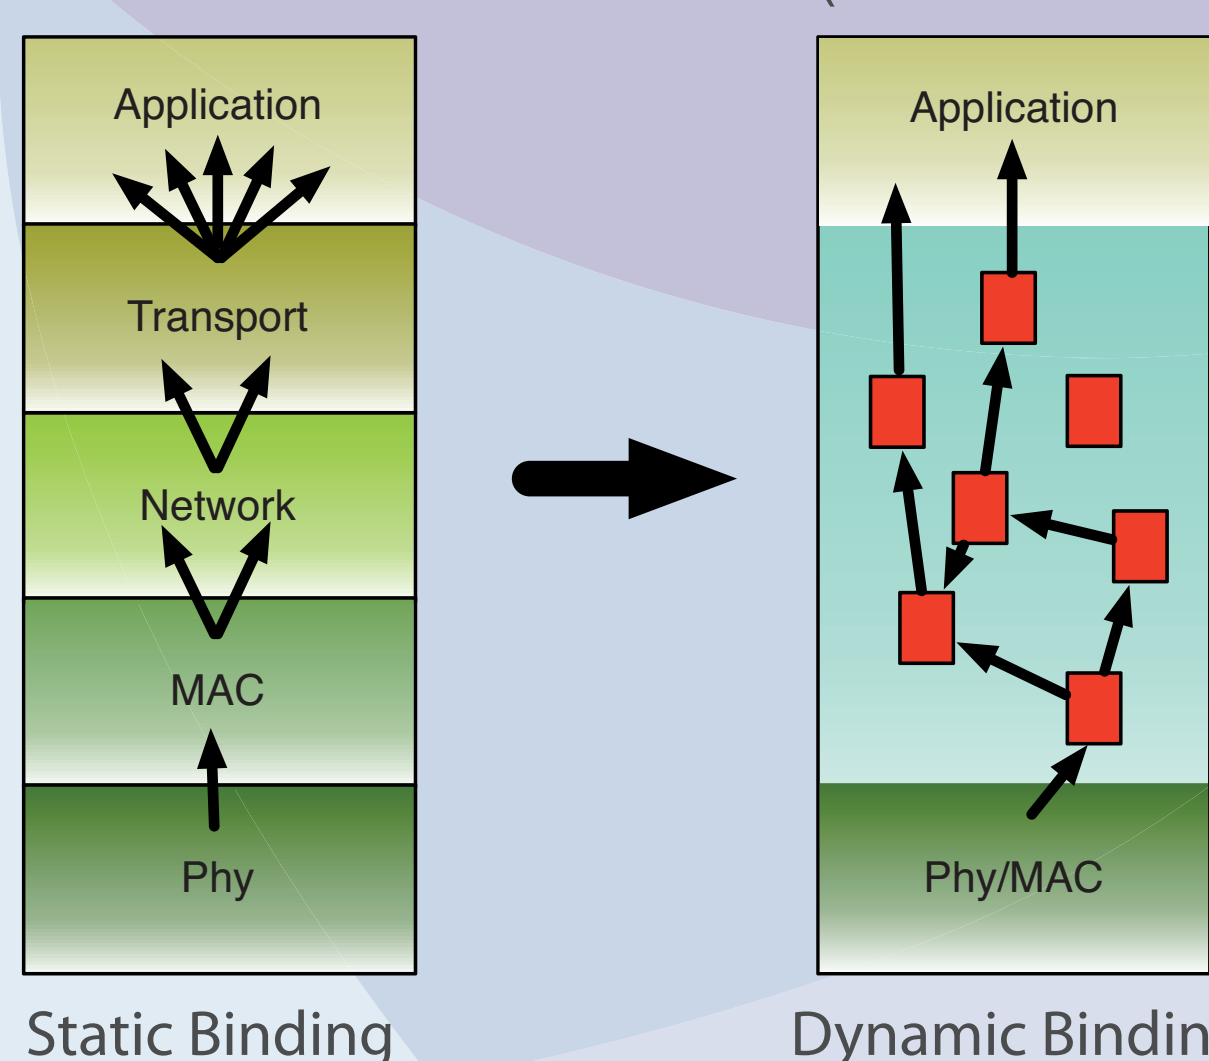




# Efficient Implementation of Dynamic Protocol Stacks

## What are Dynamic Protocol Stacks and why are they needed?

Current Internet architecture    Dynamic Protocol Stack  
(based on ANA)



### Dynamic Protocol Stacks

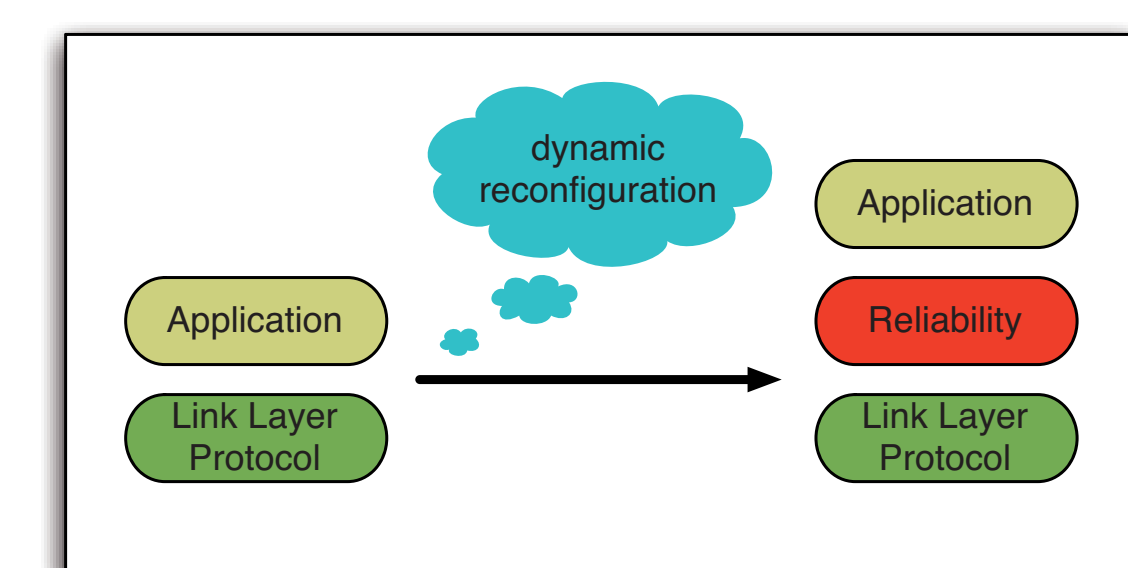
- Does not rely on TCP/IP architecture
- Does not imply any protocols to be used → meta architecture
- Divide networking functionality into functional blocks (FBs)
- Adapt protocol stack at runtime to current requirements

#### Example Scenario 1: Sensor Network

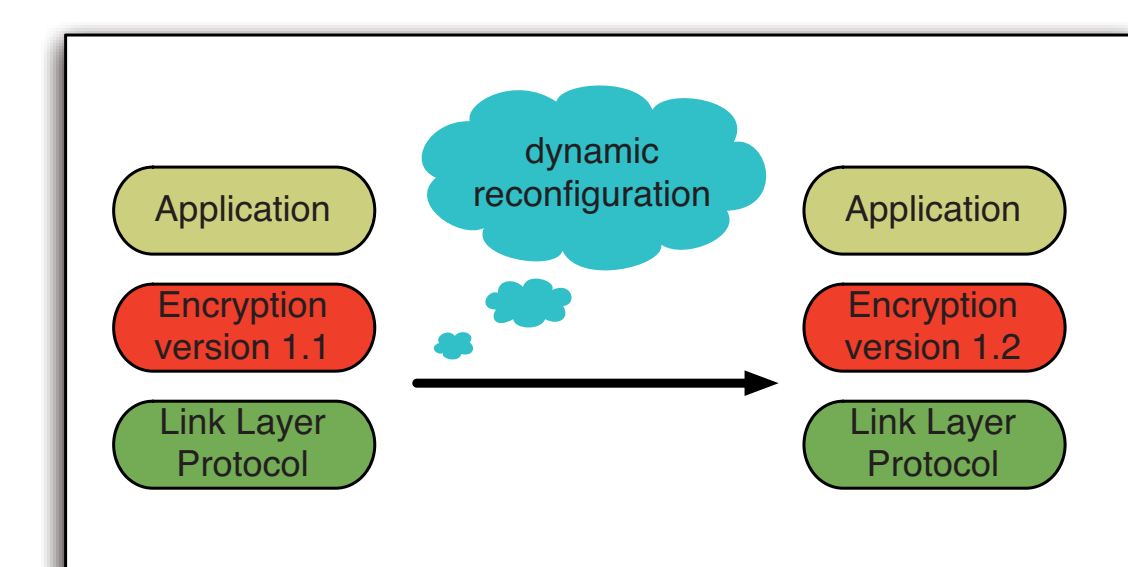
- Network conditions can change frequently
- Runtime customization of network stack to save power

#### Example Scenario 2: Video Surveillance System

- 100% uptime required
- Integration of new features and correction of bugs without service disruption



Sensor Network Scenario



Video Surveillance Scenario

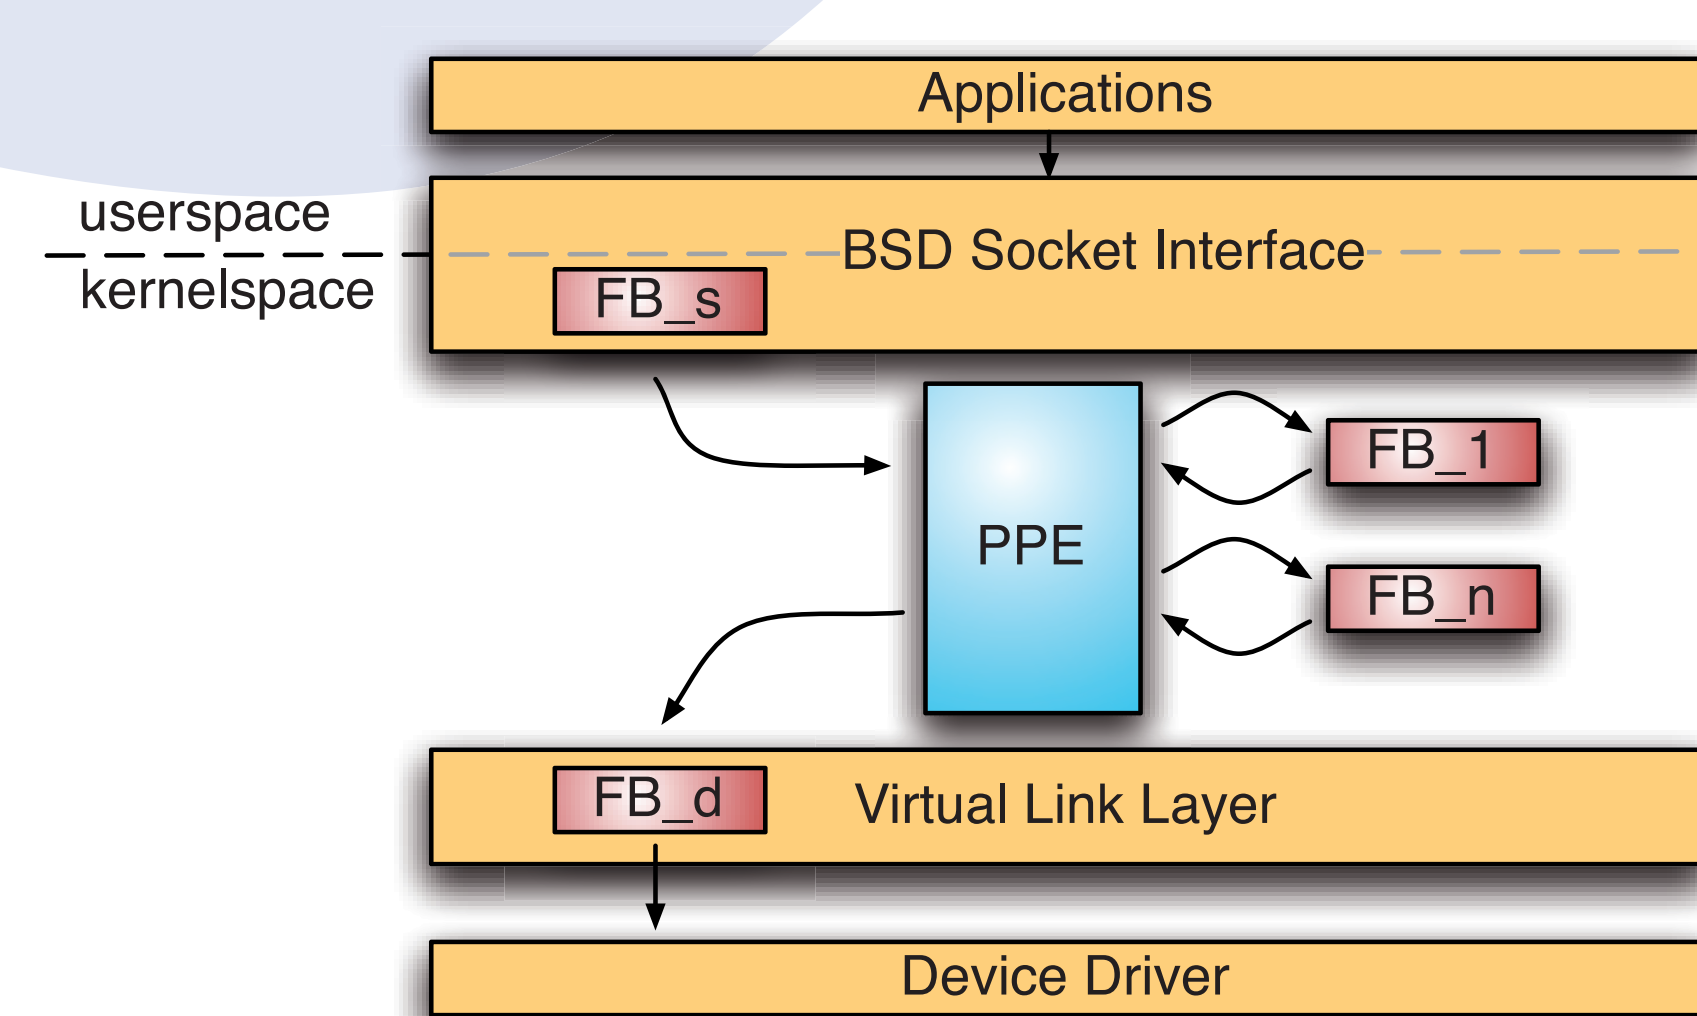
## How can Dynamic Protocol Stacks be implemented?

### Existing Approaches

- FreeBSD Netgraph: dynamic protocol configuration but only up to IP
- Click: configuration upon initialization, designed for routers, requires kernel patch
- ANA: full flexibility but only proof of concept implementation

### Our Approach

- Implementation in the Linux kernel space. Each FB is a kernel module.
- Instances of functional blocks can be generated from user space.
- Packets can be dropped, forwarded to ingress or egress, duplicated or generated.
- Packet Processing Engine (PPE) controls the execution of the FBs.
- Virtual Link Layer hides hardware and device driver peculiarities in order to allow an easy extension to other transmission channels such as Bluetooth or InfiniBand.



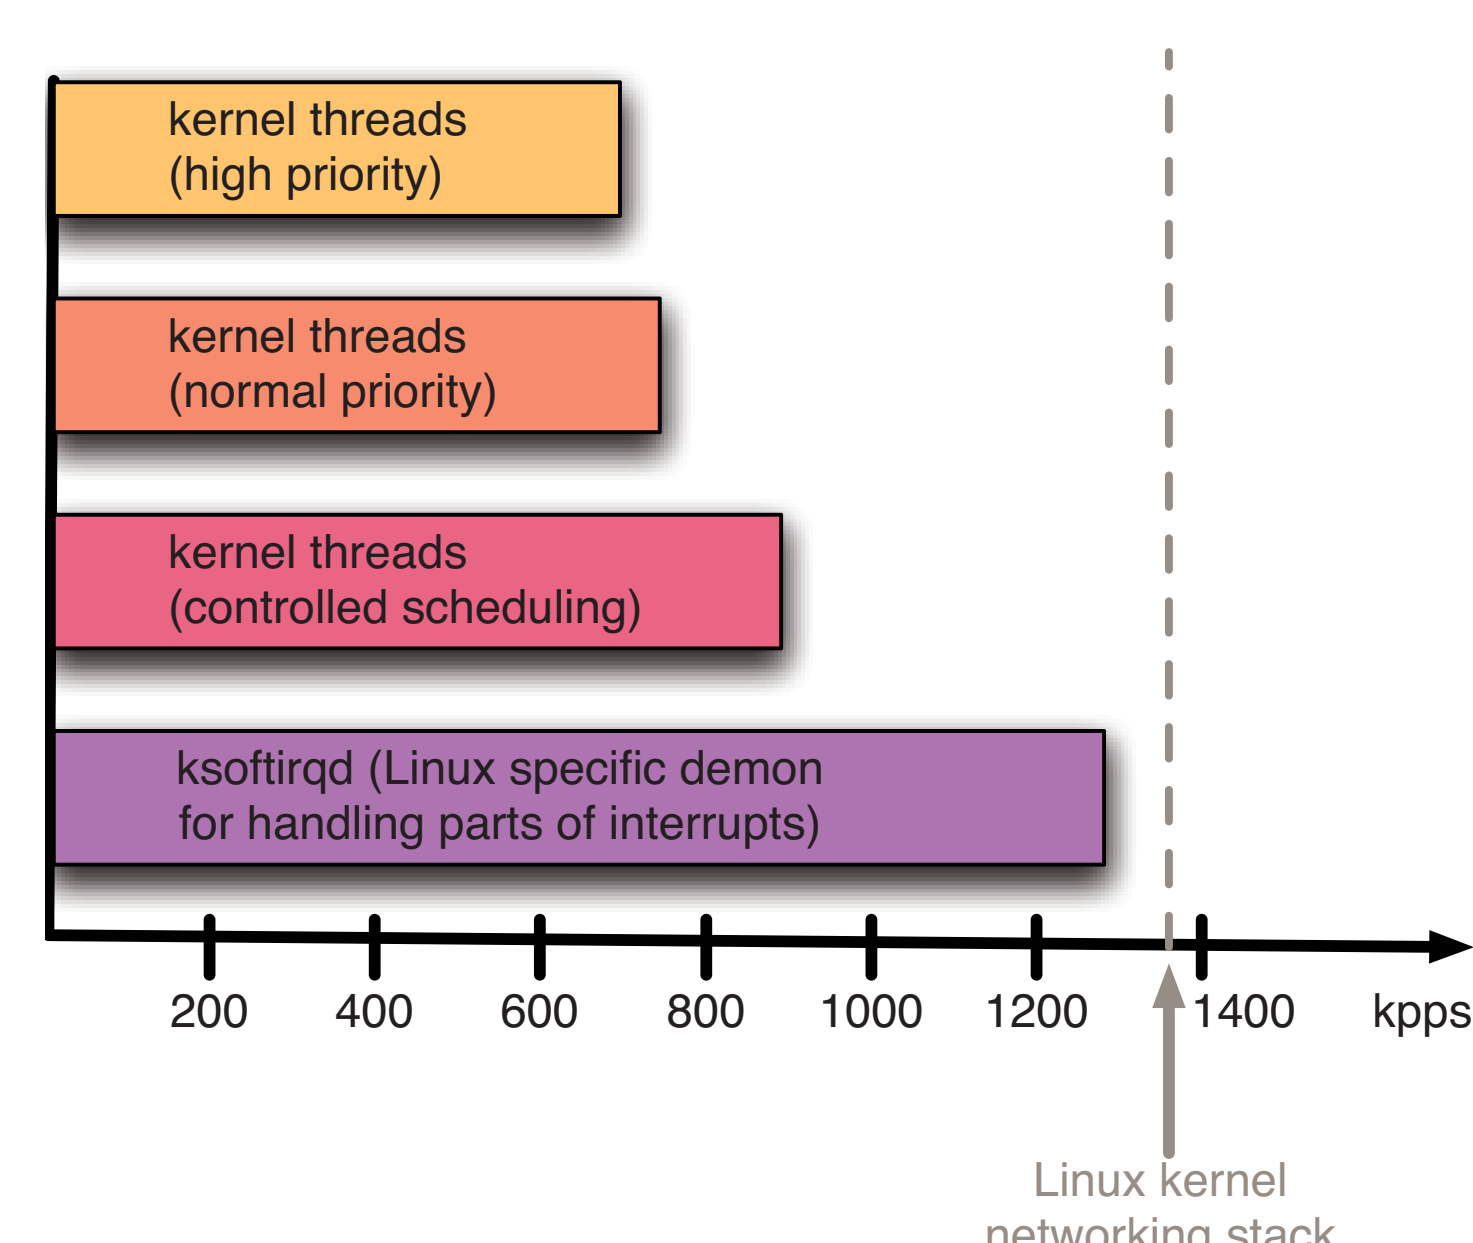
### Configuration

- `add, rm`  
Adds (removes) an FB from the list of available FBs in the kernel.
- `set`  
Sets properties of an FB with a key=value semantic.
- `bind, unbind`  
Binds (unbinds) an FB to another FB in order to be able to send messages to it.
- `replace`  
Replaces one FB with another FB. The connections between the blocks are maintained. Private data can either be transferred to the new block or dropped.

### Currently Available Blocks

- Functionality: Ethernet, Berkeley Packet Filter, BSD socket
- Packet Flow: Tee, Forwarding, Sink (packet counter)

## How fast are Dynamic Protocol Stacks?



Maximum packet reception rate for

- 64 Byte packets
- Intel core 2 Quad Q6600 2.40 GHz
- 4GB Ram
- Intel 82566DC-2 Nic
- Linux 3.0rc1

### Results

- Low level implementation influences the performance heavily.
- Interrupt processing and thread scheduling significantly influence the maximum number of packets that can be received.

## What will the future bring?

### Performance Evaluation of Real Scenarios

We will compare the performance of real scenarios in LANA with their performance in other systems (e.g., default Linux stack, Click router, etc.).

### Autonomous Configuration

We will work on mechanisms that automatically configure protocol stacks based on the applications as well as the networks needs.

### Hardware Acceleration

We plan to enhance this system with support for hardware accelerator functional blocks. The system will run on an FPGA in order to allow for the

runtime re-configuration of the hardware accelerators. The adaptation of the hardware will be done by partial re-configuration of the FPGA. For the communication between software and hardware the ReconOS operating system will be used.

### Software

The current software is available via git under the GNU General Public License from <http://repo.or.cz/w/ana-net.git>. Interested researchers are encouraged to download and try the software and report their feedback to us.

## Authors and Acknowledgements

Ariane Keller: [ariane.keller@tik.ee.ethz.ch](mailto:ariane.keller@tik.ee.ethz.ch)  
Daniel Borkmann: [dborkma@tik.ee.ethz.ch](mailto:dborkma@tik.ee.ethz.ch)  
Wolfgang Mühlbauer: [muehlbauer@tik.ee.ethz.ch](mailto:muehlbauer@tik.ee.ethz.ch)

ANCS 2011 - ACM/IEEE Symposium on Architectures  
for Networking and Communications Systems



Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

The research leading to these results has received funding from the European Union Seventh Framework Programme under grant agreement n°257906.  
[www.epics-project.eu](http://www.epics-project.eu)

